

Supervised Learning


K Nearest Neighbors

Validation


Yevgeny Seldin

Supervised Learning

- Protocol



- Example



Supervised Learning

- More examples
 - (age, gender, weight) \rightarrow height
 - (age, weight, height) \rightarrow gender
 - (height(1), height(2), height(3)) \rightarrow height(4)
- Notations
 - \mathcal{X} – sample space (e.g., $\mathcal{X} = \mathbb{R}^d$)
 - \mathcal{Y} – label space (e.g., Classification: $\mathcal{Y} = \{\pm 1\}$; Regression: $\mathcal{Y} = \mathbb{R}$)
 - $S = \{(X_1, Y_1), \dots, (X_n, Y_n)\}$ – training sample (where $X_i \in \mathcal{X}, Y_i \in \mathcal{Y}$)
 - $h: \mathcal{X} \rightarrow \mathcal{Y}$ – a prediction rule / hypothesis
 - \mathcal{H} - a set of prediction rules / a hypothesis set


K-Nearest Neighbors (K-NN)

- Algorithm: Predict X based on K nearest neighbors in S .

- Input: distance measure $d(x, x')$

- Examples:

- Euclidian distance
- Manhattan distance
- Travel distance
- Edit distance



- **The choice of d determines the success or failure of K-NN!**

Evaluation

- $\ell(y', y)$ – loss/error function
Loss for predicting y' when the reality is y
- **The loss function determines the cost of different mistakes!!!**

- Examples:

- Zero-one loss

$$\ell(y', y) = \mathbb{I}(y' \neq y)$$

$$= \begin{cases} 0, & \text{if } y' = y \\ 1, & \text{if } y' \neq y \end{cases}$$

- Squared loss

$$\ell(y', y) = (y' - y)^2$$

- Absolute loss

$$\ell(y', y) = |y' - y|$$

- Example: Fire alarm

$y' \backslash y$	no fire	fire
no fire	0	5.000.000
fire	2.000	0


“constant”

Depends
on the
house

So far

- KNN – predict based on K nearest neighbors
- Input:
 - Distance measure $d(x, x')$ – domain knowledge
- Evaluation:
 - Loss function $\ell(y', y)$ – domain knowledge


How to pick K ?




- What is good/bad about small K ?
 - Say, $K=1$?
- What is good/bad about large K ?
 - Say, $K=n$?

How to pick K ?

- Target: minimize the expected loss
 - $L(h_{KNN}) = \mathbb{E}[\ell(h_{KNN}(X), Y)]$
- Assumption
 - (X, Y) are sampled from a fixed (unknown) distribution $p(X, Y)$
 - The expectation is with respect to $p(X, Y)$
- Challenge: $p(X, Y)$ is unknown, and so is $L(h_{KNN})$
- How to estimate $L(h_{KNN})$?
 - Use the empirical loss $\hat{L}(h_{KNN}, S) = \frac{1}{n} \sum_{i=1}^n \ell(h_{KNN}(X_i), Y_i)$
 - What is $\hat{L}(h_{1NN}, S)$?
 - In general, $\hat{L}(h_{KNN}, S)$ is an underestimate of $L(h_{KNN})$.



Validation




- **Assumptions**

- $\{(X_1, Y_1), \dots, (X_m, Y_m)\}$ are independent identically distributed (i.i.d.)
- And come from the same distribution as new samples (X, Y)

- $\hat{L}(h_{KNNS_t}, S_{val})$ is an **unbiased** estimate of $L(h_{KNNS_t})$

- $\mathbb{E}[\hat{L}(h_{KNNS_t}, S_{val})] = \mathbb{E}\left[\frac{1}{m} \sum_{i=1}^m \ell(h_{KNNS_t}(X_i), Y_i)\right] = \frac{1}{m} \sum_{i=1}^m \mathbb{E}[\ell(h_{KNNS_t}(X_i), Y_i)] = L(h_{KNNS_t})$
- From the perspective of h_{KNNS_t} the samples in S_{val} are indistinguishable from new samples (X, Y)

Selection of K




- **Selection introduces bias!**


- Each $\hat{L}(h_{KNN}, S_{val})$ is an unbiased estimate of $L(h_{KNN})$
- But $\hat{L}(h_{K^*NN}, S_{val})$ is a **biased** estimate of $L(h_{K^*NN})$!!!
 - From the perspective of h_{K^*NN} the samples in S_{val} are distinguishable from new samples (X, Y)
 - $\mathbb{E}[\ell(h_{K^*NN}(X_i), Y_i)] \neq \mathbb{E}[\ell(h_{K^*NN}(X), Y)]$

dependent!
in S_{val}

Illustration of Selection Bias




A bag of coins with bias p

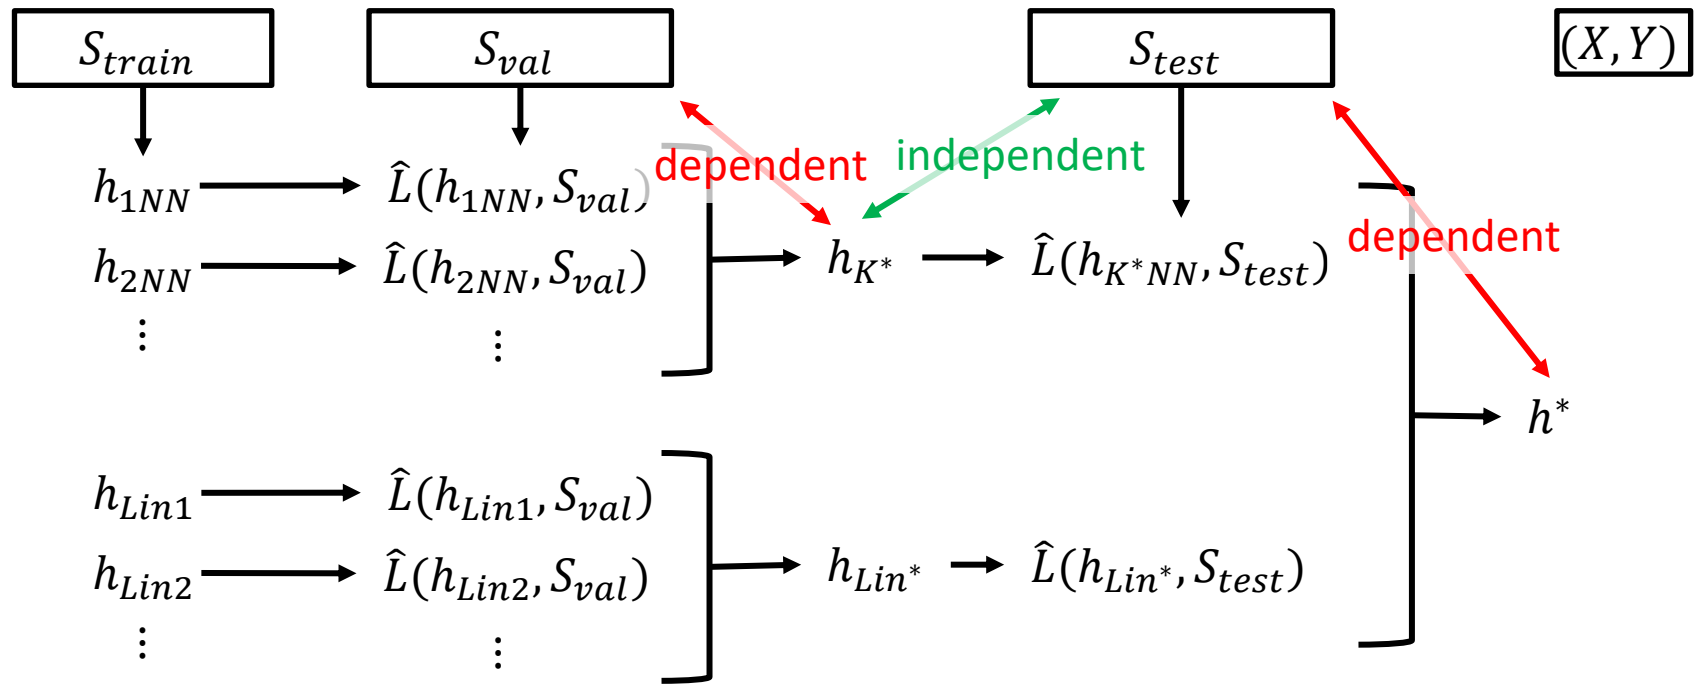


- $i^* = \arg \min_i \hat{p}_i$
- While each \hat{p}_i is an unbiased estimate of p : $\mathbb{E}[\hat{p}_i] = p$
- \hat{p}_{i^*} is a **biased** estimate of p : $\mathbb{E}[\hat{p}_{i^*}] \neq p$
- Outcome-based selection introduces bias!

So how can we estimate $L(h_{K^*_{NN}})$?

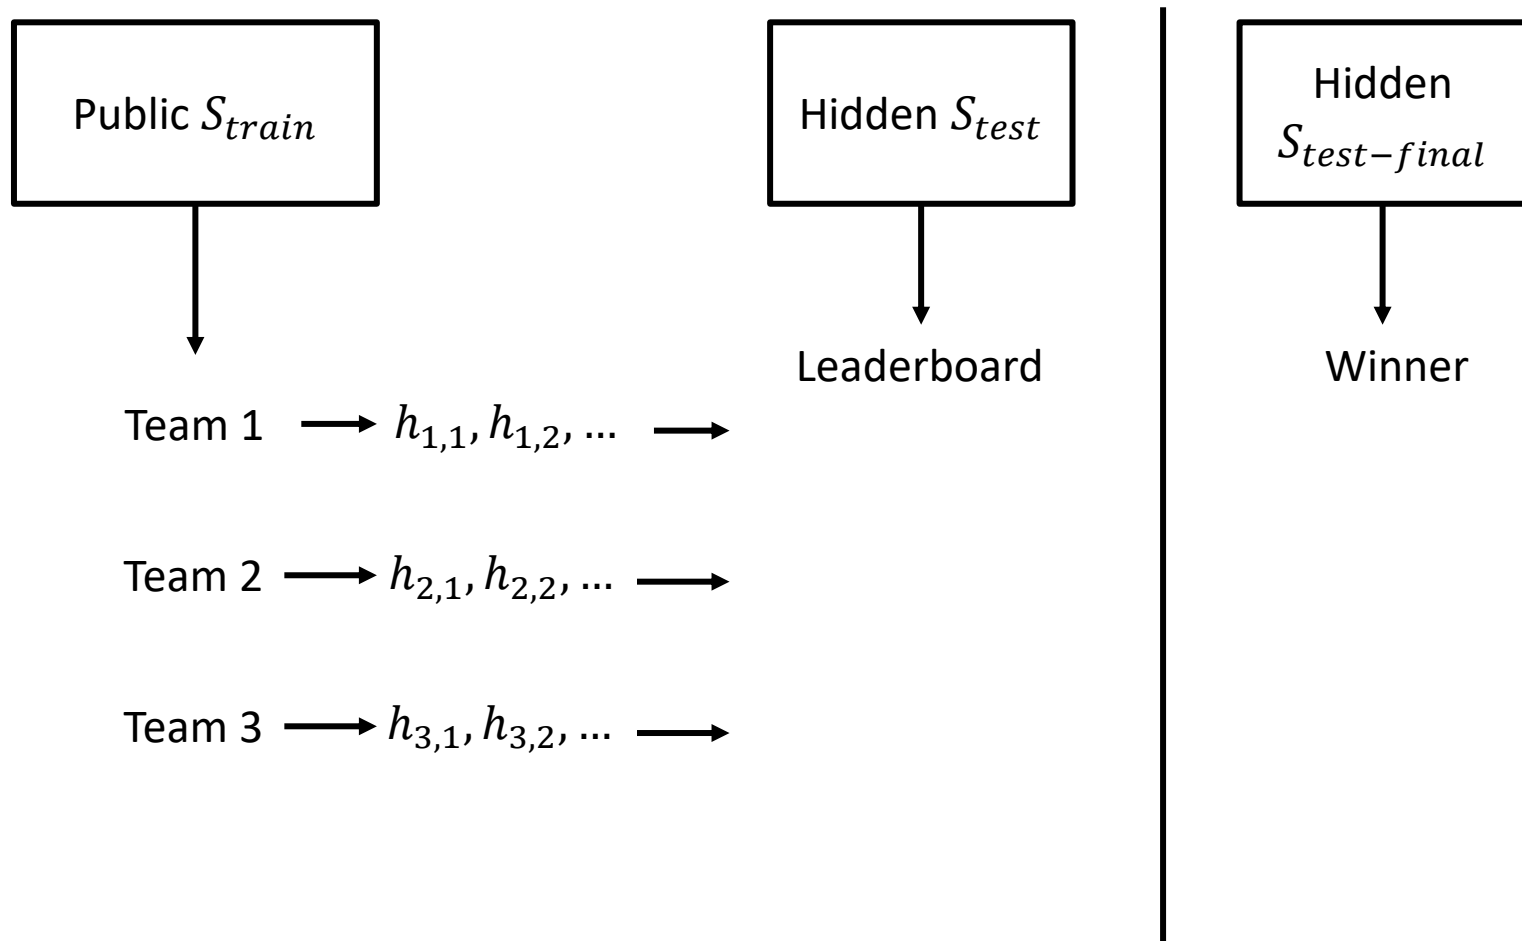


Testing



- It's not about how you call it; it's about how you use it!!!
- $\hat{L}(h^*, S_{test})$ is a **biased** estimate of $L(h^*)$!

Respectable ML Competitions



Division of Responsibilities

- Evaluation/error measure $\ell(y', y)$ – domain knowledge
- Distance measure $d(x, x')$ – domain knowledge
 - But validation can be used to select a distance measure out of a set of candidates
- Selection of K and d – validation
- How to split the data (select m) – coming next

How to split the data into train/test/...?

- Consider the following extremes:



$$m = 1$$



$$m = n - 1$$

- $m = 1$
 - $\hat{L}(h, S_{test}) \in \{0,1\}$, never approaches $L(h)$
 - $m = n - 1$
 - The training procedure only observes one label
- Do we need to reshuffle the data before splitting?
 - Theory: the data are assumed to be i.i.d., so it does not matter
 - Practice:
 - Yes, if the data are sorted by irrelevant parameter
 - No, if data order carries information relevant for testing, e.g., ordering by time

What can be said about $L(h)$ based on $\hat{L}(h, S_{val})$?

- $\hat{L}(h, S_{val})$ is an unbiased estimate of $L(h)$
- But consider the case $m = 1$:
 - $\hat{L}(h, S_{val}) \in \{0,1\}$ – never close to $L(h)$!
- Being unbiased is neither sufficient, nor necessary
- We need concentration!

Relation to “coin flips”

- $Z_i = \ell(h(X_i), Y_i) \in \{0,1\}$
 - Bernoulli random variable, “a coin flip”
- $\mathbb{E}[Z_i] = \mathbb{E}[\ell(h(X_i), Y_i)] = L(h) = p$
 - The bias of the coin
- $\hat{L}(h, S_{val}) = \frac{1}{m} \sum_i^m Z_i = \hat{p}_m$
 - An average of m “coin flips”
- How far can \hat{p}_m be from p ?